

Token generation latency vs. prefix length

